



Build systems

To a simple solution
To a complicated problem

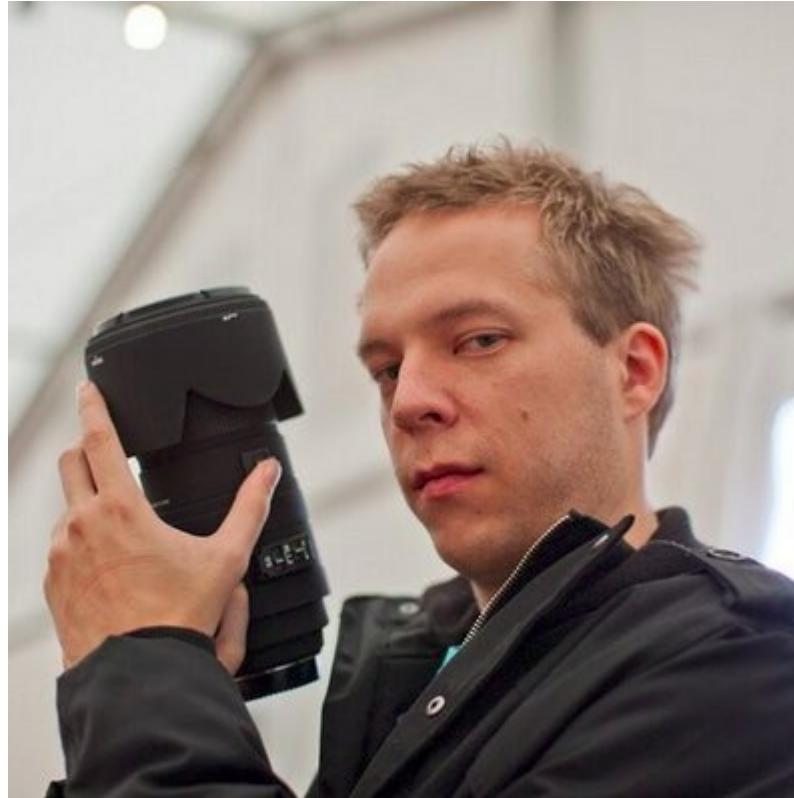


**Who are we
doing this for**

The creator of autotools?



Build-tool specific experts?

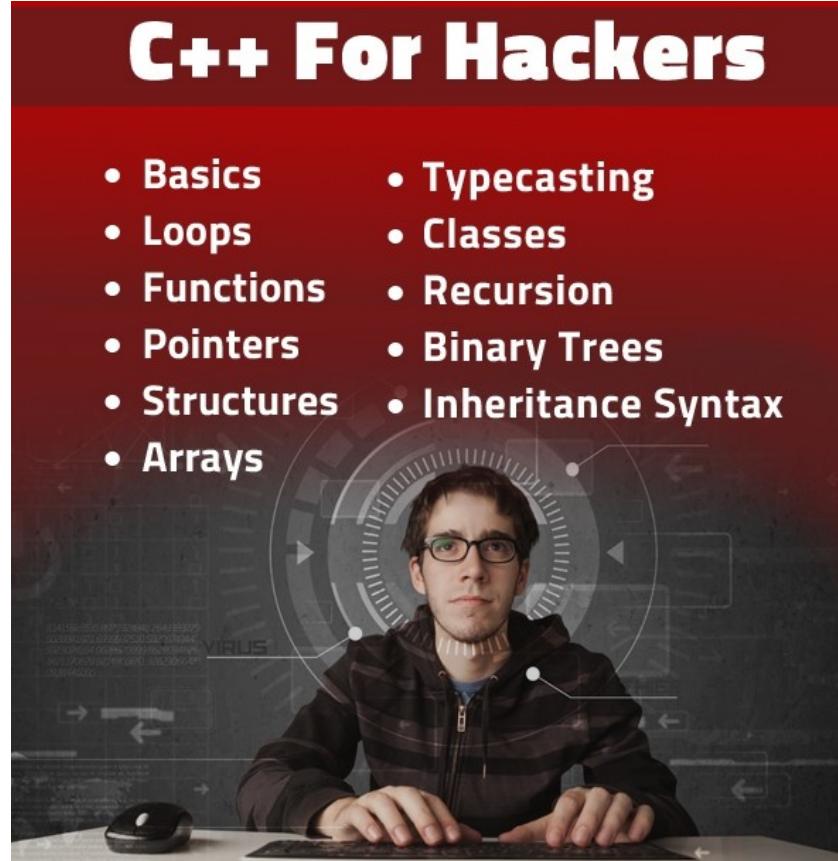


Source:
CppCast,
Jussi Pakkanen,
Meson author

People that attend conferences?



Familiar with the language?



Source:
[http://www.hackinsight.org/
free-content,33.html](http://www.hackinsight.org/free-content,33.html)

New to the language



New to programming entirely



Students seeing a compiler for the first time in their life



Children learning new skills





TOMTOM 

© 2018, Peter Bindels

My target

- You can read and write
- You can understand English keywords
 - Somewhat
- You have a computer
 - But a 10-year old computer should suffice

My target

- Make it really simple to start
- Allow the users to learn beyond
- Offer a cheap and smooth upgrade path

Simplicity

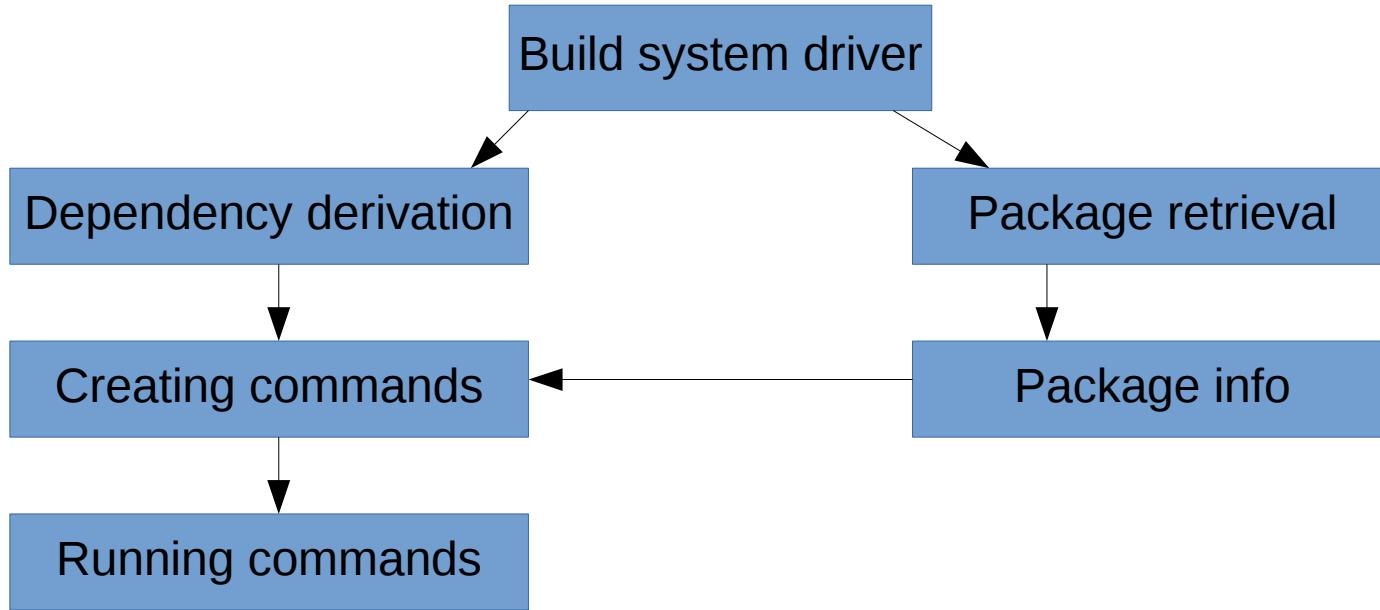
- You saw Kate's talk yesterday

“Is this really all I need to change?”

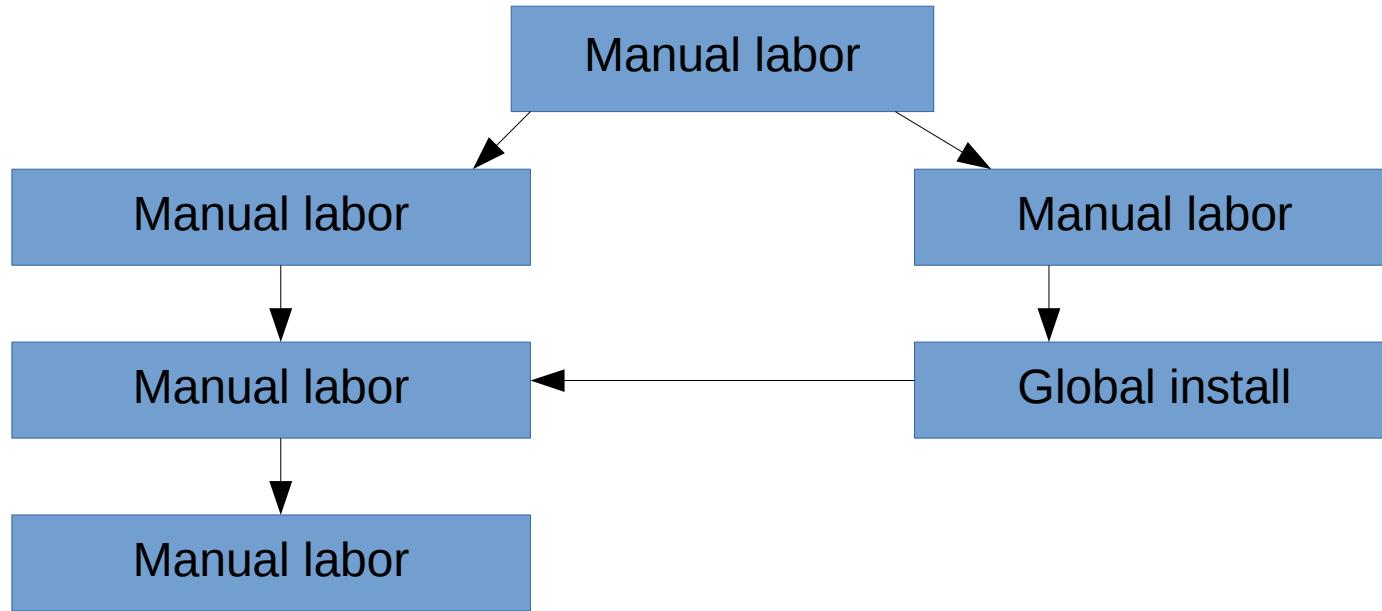
THE #1 PROGRAMMER EXCUSE FOR LEGITIMATELY SLACKING OFF: "MY CODE'S COMPILING."



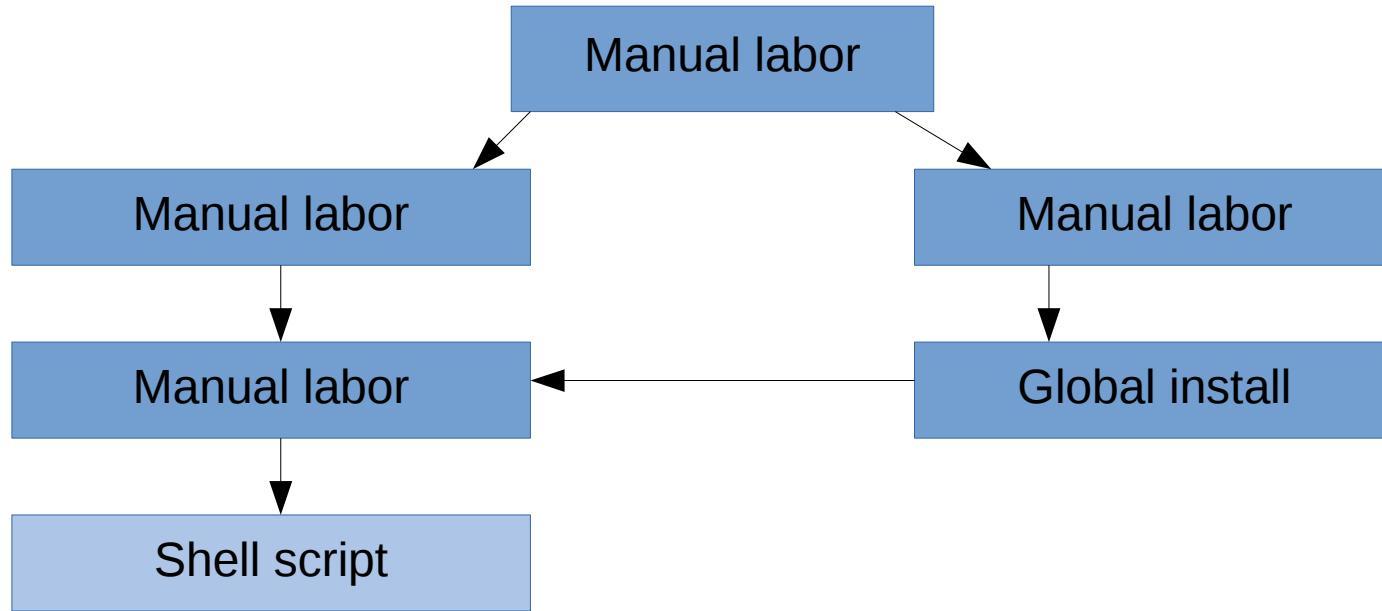
The parts of a build



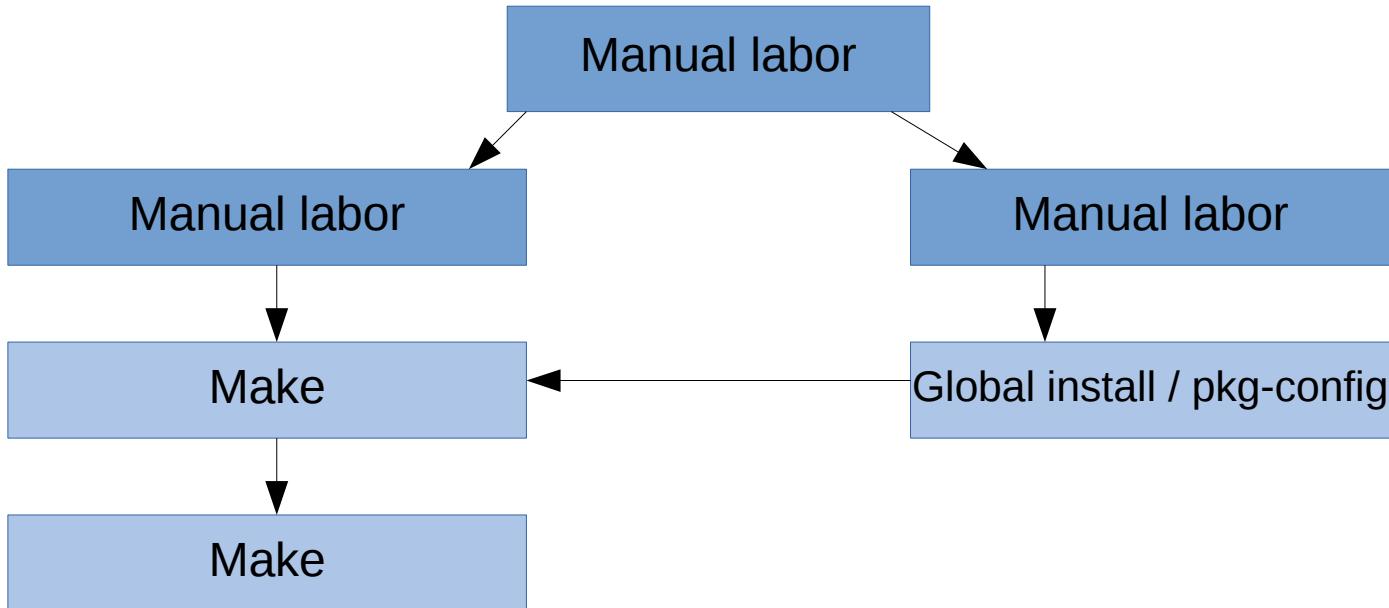
Way before time began



Way before time began



The Make era



HOW STANDARDS PROLIFERATE:

(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC)

SITUATION:
THERE ARE
14 COMPETING
STANDARDS.

14?! RIDICULOUS!
WE NEED TO DEVELOP
ONE UNIVERSAL STANDARD
THAT COVERS EVERYONE'S
USE CASES.



Soon:

SITUATION:
THERE ARE
15 COMPETING
STANDARDS.

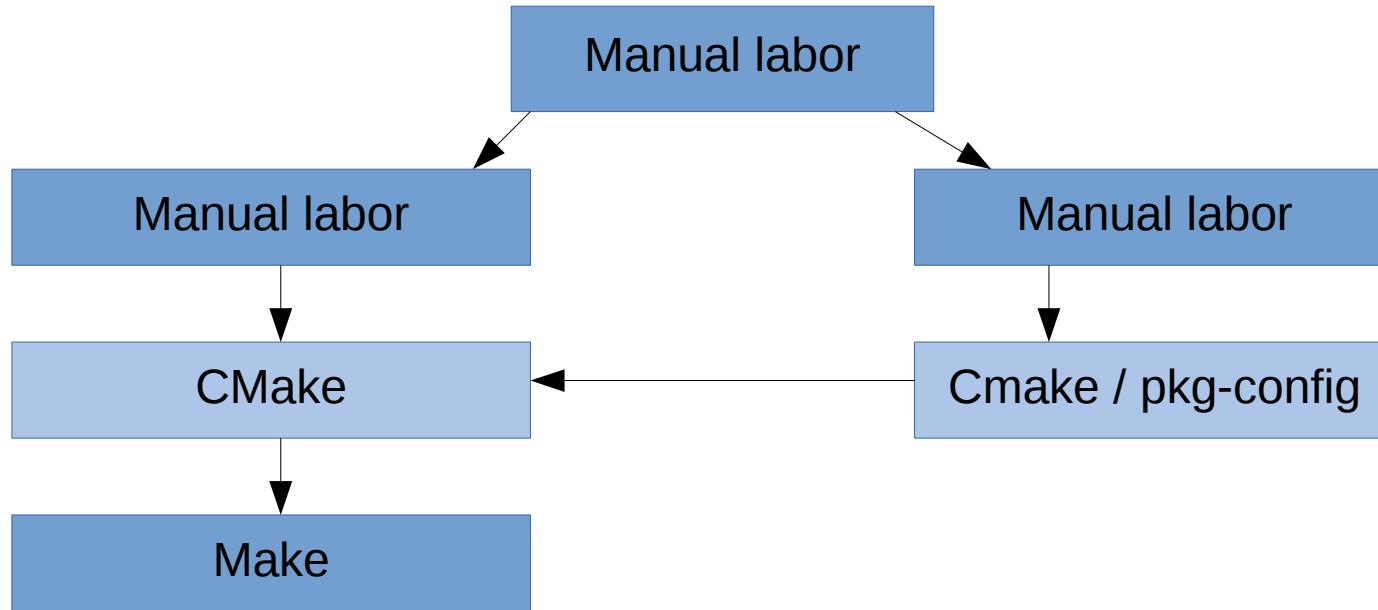
All the makes

- bmake
- cmake
- dmake
- emake
- fmake
- gmake
- hmake
- imake
- jmake
- kmake
- mmake
- nmake
- omake
- pmake
- qmake
- rmake
- smake
- tmake
- umake
- vmake
- wmake
- xmake
- ymake
- zmake

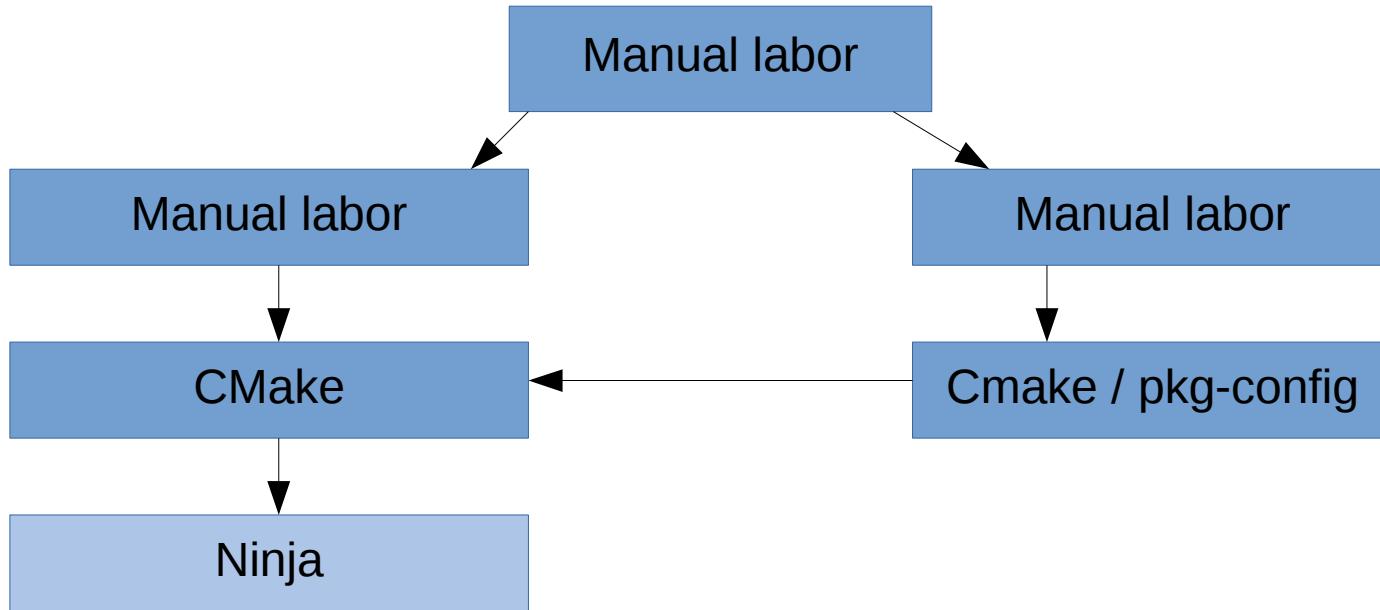
If you want to make another make

- amake
- lmake
 - Not to be confused with imake, of course
- But my list is from 2011, so maybe not
- What about ☺make ?

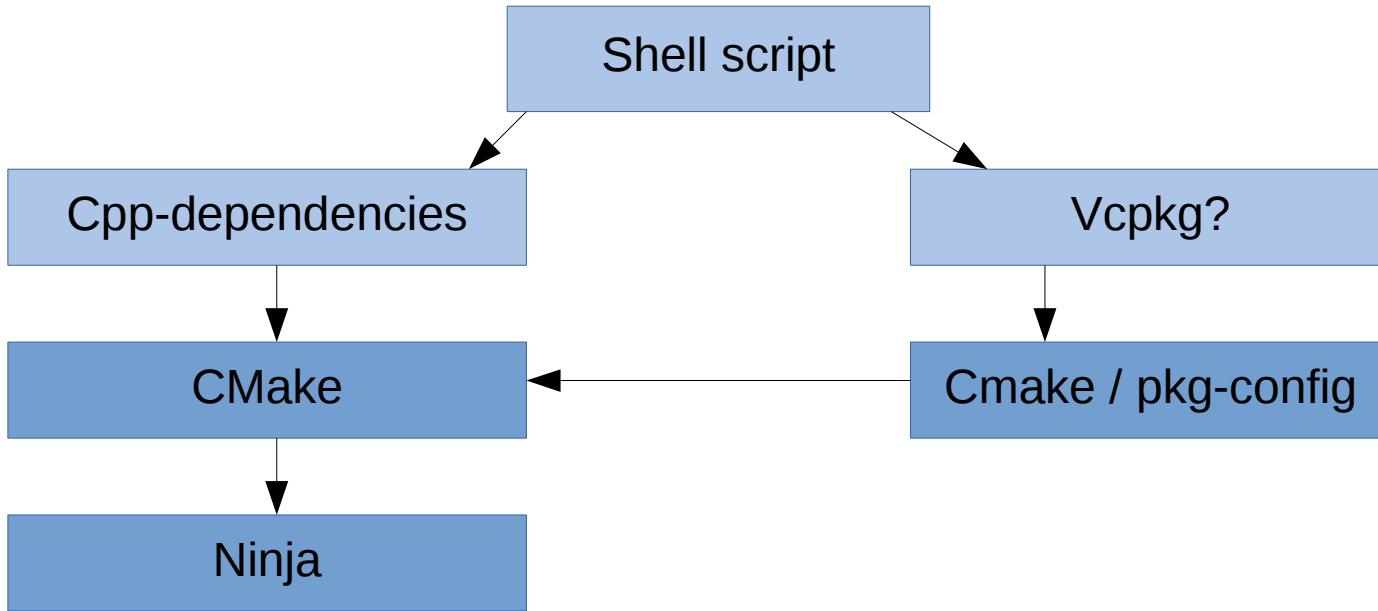
The Cmake era

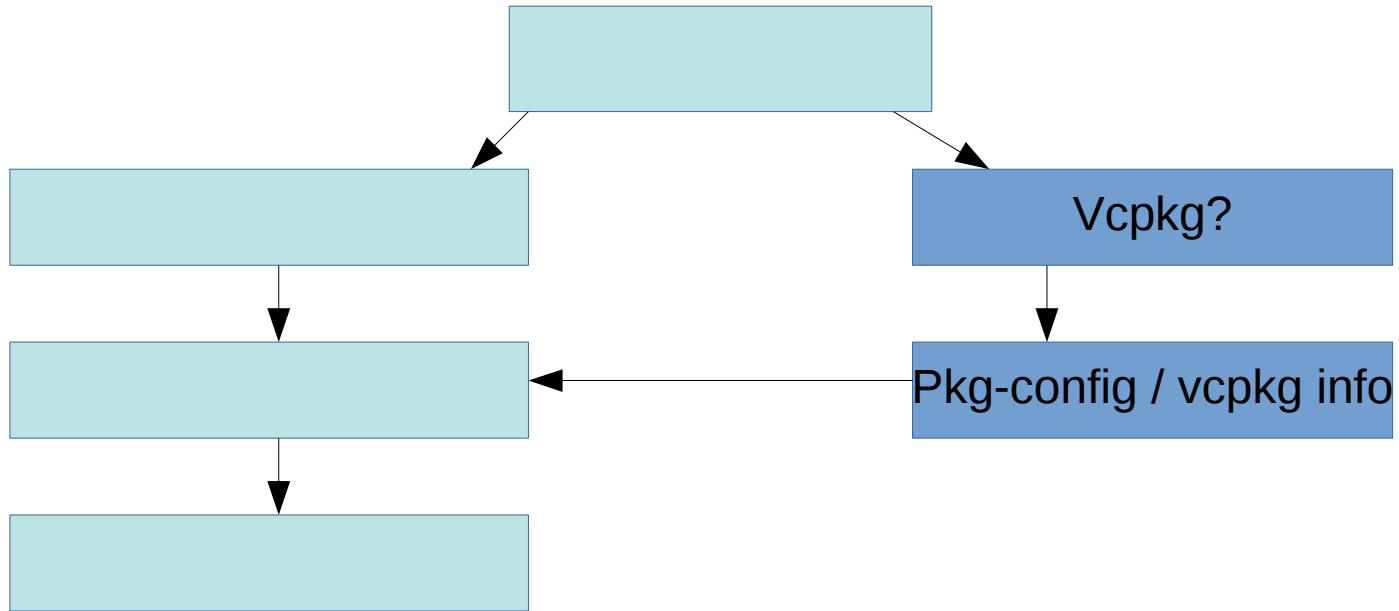


Beyond make



Recent history





Why do build systems always tend to complexity



All the corner cases



Total build sequence

- Build object files
- Pack static libraries
- Link executables
- Link shared libraries
- Run unit tests
- Run code formatter

Total build sequence

- Build object files
- Pack static libraries
- Link executables
- Link shared libraries
- Run unit tests
- Run code formatter
- Run static code analysis

Total build sequence

- Build object files
- Pack static libraries
- Link executables
- Link shared libraries
- Run unit tests
- Run code formatters
- Run static code analysis
- Upload to embedded target
- Send to binary package repository
- Download code from git
- Update version file & commit
- Git submodule init
- update
- Create release notes
- Announce on twitter
- Wget script | sudo bash -c

And more!

Modules

What are modules

- From build system point of view

Generated code

- Generated source files
- Having your source generator part of the same build
- That depends on some other generated file itself
- Recursively repeated.

Build system types

- Procedural
- Functional
- Declarative / descriptive
- Prescriptive / Generative

Procedural

- Your build steps
- In the order
- To run them

Procedural

- Simple to grasp
- Can be reliable

Procedural

- 100% manual scripting
 - Omissions
 - Bugs
 - Typos
 - Unintended deviations

Procedural

- Requires full knowledge of compile process
- No parallelism
- No incremental compilation
- **Very** easy to mess up

Functional

- Define how to derive the build steps
- Let the build system figure out what follows
- Excellent if compile steps and tools vary a lot

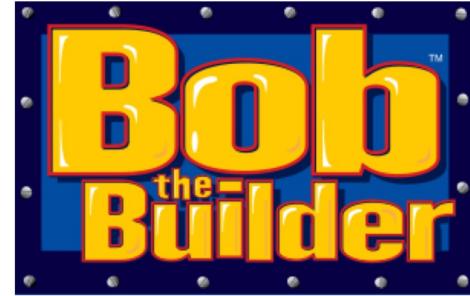
Functional

- Handles incremental builds
- Can be very reliable
 - But make doesn't warn if you mess it up
- Portable
 - (arguably)

Functional

- Nowadays very unpopular
- Really hard to diagnose and fix
- Implementation incompatibility





Functional

- No good solution for dependencies
- Minor historical limitations

Declarative / Descriptive

- Most build systems
- About as verbose as the build script
- Allows arbitrary scripting

Who here has seen build scripts

- Special cases in some targets
- Custom hacks
- // TODO: remove this in 5 years
- // 2005-06-28 Temporary hack

Who knows *FOR SURE* that you don't have this?

Remember that you can override - for example - add_library in cmake to do something else entirely.

This could be in somebody else's scripts that you include

Prescriptive/Generative

- Instead of writing down what to build, derive what to build
- Allow minimal room for adjustments and exceptions
 - Nobody else can either

Prescriptive/Generative

- Complicated build tool
- Requires knowledge of what the tool expects

So how do we break the cycle of complexity?



Explicitly target beginners

- Very little configuration required
 - To compile your Hello World it should require absolutely no configuration.
 - As far as is reasonably possible, it should require no configuration.
 - If there is configurability, it should be restricted to acceptable minor adjustments.
 - Name and type of component

Explicitly target beginners

- Limit creativity
 - Make it *not* turing complete.

Explicitly target beginners

- Clearly delineate what the tool will accomodate, and what it will not.
- Limit deployment type proliferation
- Error for build-breaking conditions
- Warn for things that cause build-breaking conditions

Make it prescriptive
Make it generative

Why do prescriptive systems fail

- People hate changing habits
- People dislike being prescribed things



Works for Java



Works for Basic



***** COMMODORE 64 BASIC V2 *****
K RAM SYSTEM 38911 BASIC BYTES F



© 2018, Peter Bindels

Works for Pascal



Pascal

***** COMMODORE 64 BASIC V2 *****
K RAM SYSTEM 38911 BASIC BYTES F



© 2018, Peter Bindels

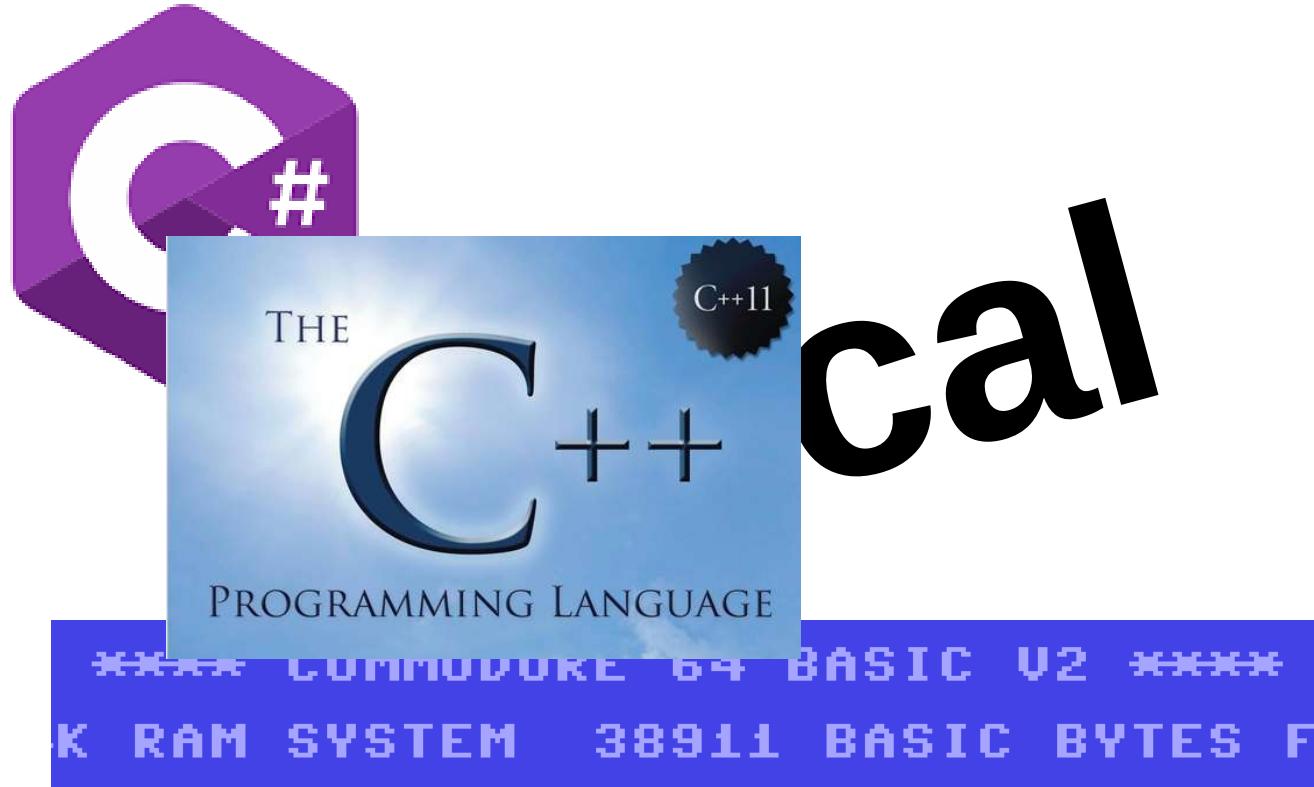
Works for C#



Pascal

***** COMMODORE 64 BASIC V2 *****
K RAM SYSTEM 38911 BASIC BYTES F

Why not C++?



Inheritance from C

- Started before build systems
- Did not standardize one
- People forked and innovated
 - => zoo of tools & inscrutable build solutions
- C++ just went with this

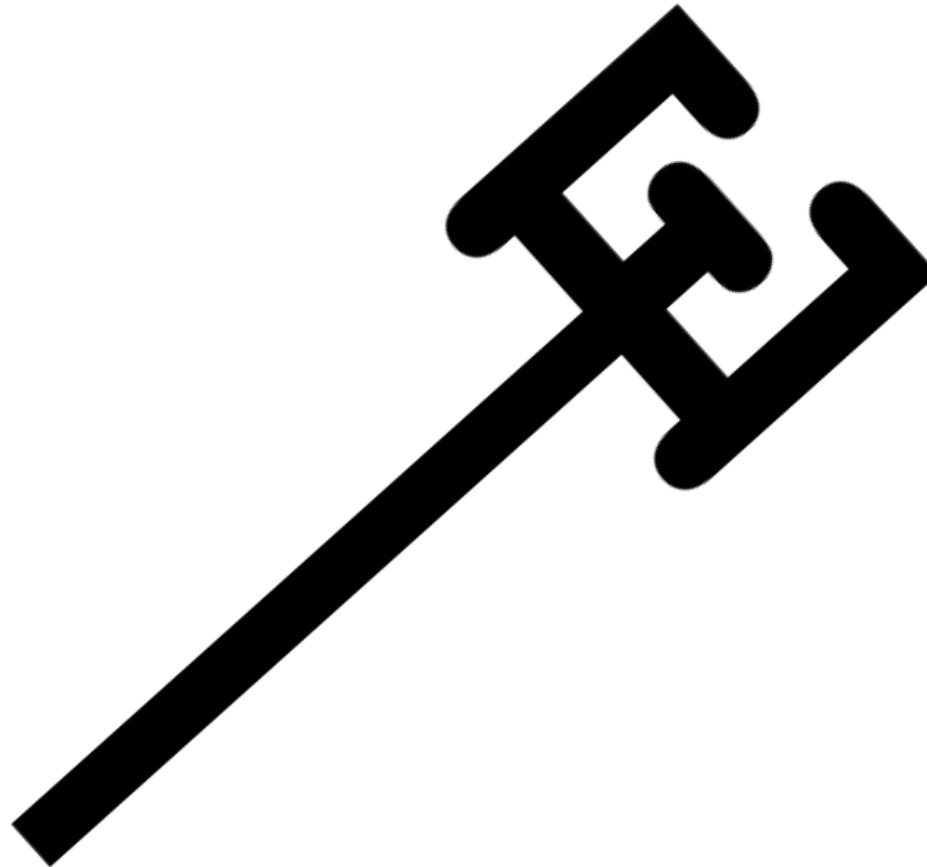
How do you make a
prescriptive solution
work?

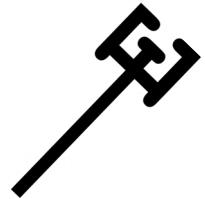
Prescribe the status quo

- Prescribe what we already do
 - in a structured way
 - Make everything derivable from the code
 - Allow the user to override
 - but only where we must allow variation

Standardized project layout

- Per component
 - src
 - include
 - test
 - doc
 - examples
 - ...





Pitchfork proposal by Colby Pike

- <https://ol.reddit.com/r/cpp/comments/996q8o>
- Take the current defacto standard
- Codify it
- Tweak it where you must
- <https://github.com/vector-of-bool/pitchfork>

What this offers to tools

- Warning people for relying on implementation details
- Implicit "at home" feeling in a new code base
- Automatically detecting dependencies
 - Zero-configuration project layout information

Dependencies

- Compile time
- Link time
- Run time

Compile time

- `#include <pixel.hpp>`
- `import std.io;`

Link time

- extern double d;
- void f();

Runtime

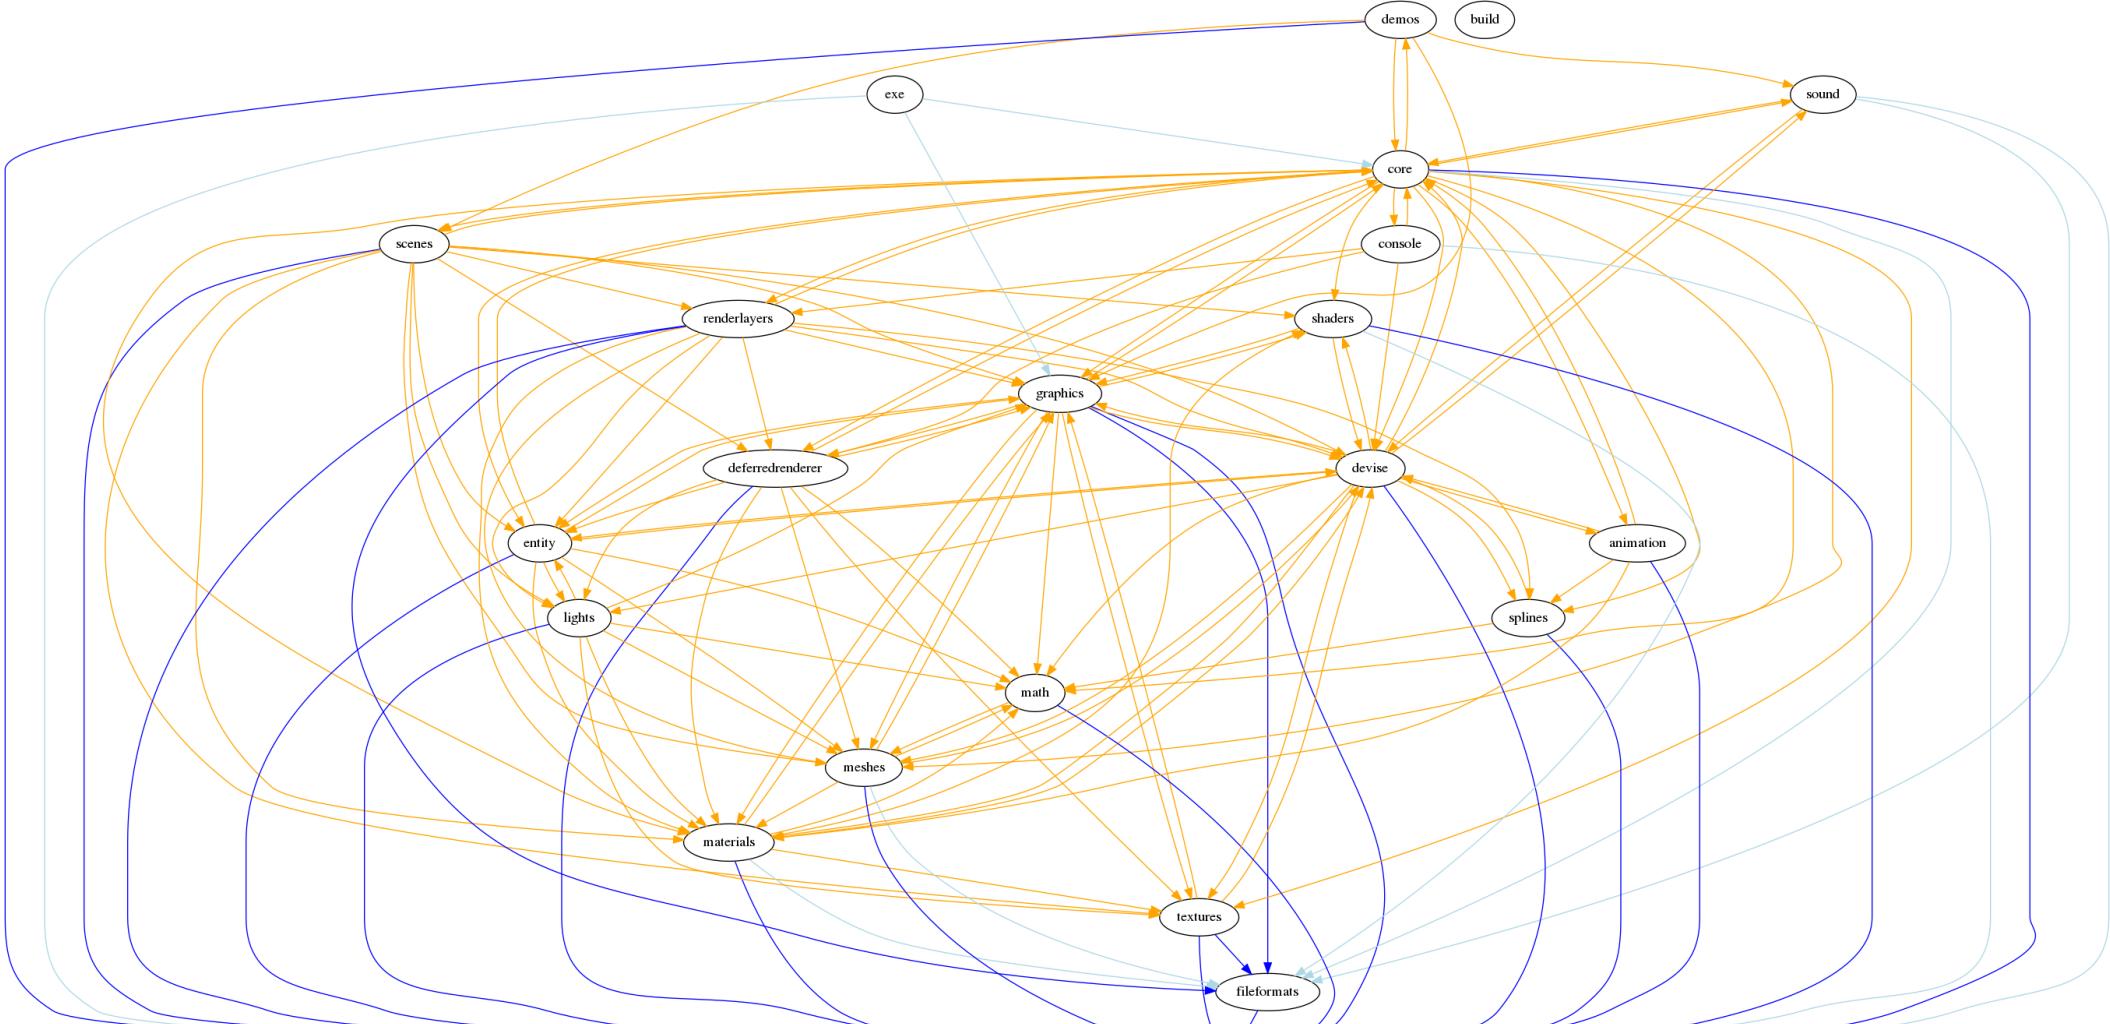
- Shared libraries
- Plugins
- Database links
- Operating system functions
- External hardware
- ...

Dependencies

- Compile time is better than link time
 - Detectable dependencies are good
- Link time is better than run time
 - Breaking a build beats having complaining customers

Manually specified

- target_link_libraries
- Relative or absolute includes work around this
- Overspecified
- Underspecified





Humans

- Humans
 - Smart
 - Fallible
 - Computers
 - consistently dumb
 - very consistent
- (error rate >2%)
- (error rate < 0.0000000001%)

**DRY
DON'T
REPEAT
YOURSELF**



Write Everything Twice

Not hardcoding inferred data

==

it's up to date

Cpp-dependencies

- Read all your code
 - Extract includes
 - Extract imports
- Derive what depends on what

Cpp-dependencies

- Discover transitive dependencies
- We know included files
 - Who includes them
 - What include path they need

Cpp-dependencies

- Component-wise thinking
 -  Gives us component roots
 - CMakeLists.txt location also works

Cpp-dependencies

- We know component locations
- We know all constituent files
- We know public and private dependencies
- We know public and private include paths
- This is 95% of what your build needs...

Cpp-dependencies at TomTom

- TomTom proprietary code base
 - 7M LoC
 - 500 components
 - Lot of test coverage
- 2012 4-core laptop with sufficient memory
- Expectancy: Should be able to match C preprocessor in speed

Cpp-dependencies at TomTom

- 2s run time with hot cache
- SSD speed limited with cold cache
- Interactive analysis
- Cycle detection & prevention
- Autogenerating cmakelists
 - Last time I checked we did about 60%

Cpp-dependencies at TomTom

- MeetingC++ talks in 2016 and 2017
- So how to make this into a build system?

Evoke

- A simple solution for the build system problem
- Evoke on Merriam-Webster:
 - 1 : to call forth (or up)
 - 2 : to re-create imaginatively

Basic idea

- Take
 - Cpp-dependencies
 - Cmake with fully generated build files
 - Ninja to build something
 - (Package manager to find needed packages)
- Roll that into one

Later goals

- Run continuously
- Cross-compile for other platforms

Bugs in all build systems now

- Race conditions
 - Save a file while it's being used for a compilation
 - the compile timestamp for object file will be **after** source file
 - Won't trigger rebuild if it was successful
 - Will cause very hairy bugs if you continue
 - Solution: When done with compiling, pre-date the file to the start of compilation

Bugs in all build systems now

- Incremental builds hide warnings
 - Files with warnings build successfully
 - Incremental build skips them
 - Warnings are hidden
- This is why -Werror is popular.
- Solution: Cache warnings & re-emit if not rebuilding
 - Breaks time-dependent commands

Design

- Take all includes and imports, derive file-level dependencies.
 - Ask package manager for packages corresponding to missing files?
 - If the package manager did something, rebuild the tree including those packages.

Design

- Translate file-level dependencies to component-level dependencies (like cpp-dependencies)
- Translate the components and dependencies into command lines for a given target with file-level dependencies (like CMake)
- Invoke the low-level build runner to handle these (like Ninja)
 - Is currently built-in, but could be externalized to Ninja, Make or something like it

Current limitations

- Derived dependencies from includes and imports
- Imperfect parser
 - >99.99% recognized though
- There is the 5% that cannot be done.

The ifdef problem

```
#ifdef X  
  
#include "file1.h"  
  
#else  
  
#include "file2.h"  
  
#endif
```

Current limitations

- Toolchain switching (Clang)
- Modules is completely untested.
 - Implemented as per current TS
 - Current Clang matches the design; I just haven't gotten around to it.
 - Does not include this weeks' changes

Current limitations

- Untested implementation
 - Android
- Unimplemented
 - OSX
 - Windows
 - iOS
 - No reason they can't be, they just aren't now

Current limitations

- Package manager links
 - Dependency import is not implemented properly
 - Demo uses a tiny stub. The stub conforms to the interface.
 - Tested with a self-written package system
 - Miles from release quality

- Open to contributors
 - Come talk to me after the talk
 - Reach out
 - Even – *especially* – if you can't code
-
- <https://github.com/dascandy/evoke>

Demo

If we have time

Hello World from scratch

Pixel showcase with Evoke

Let's make the C++ world a
simpler
place

- Peter Bindels
 - TomTom Intl
 - @dascandy42
 - dascandy at discord&slack
-
- <https://github.com/tomtom-international/cpp-dependencies>
 - <https://github.com/dascandy/evoke>